

```
In [1]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression
from sklearn.svm import LinearSVC
from sklearn.metrics import accuracy_score, classification_report
```

```
In [2]: df = pd.read_csv("spam.csv", encoding="latin-1")
df = df[['v1', 'v2']]
df.columns = ['label', 'message']
```

```
In [3]: df.head()
```

```
Out[3]:
```

	label	message
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

```
In [4]: df['label'] = df['label'].map({'ham': 0, 'spam': 1})
```

```
In [5]: X_train, X_test, y_train, y_test = train_test_split( df['message'], df['label'], test_size=0.2, random_state=42)
```

```
In [6]: # TF-IDF Vectorization
vectorizer = TfidfVectorizer(stop_words='english')
X_train_tfidf = vectorizer.fit_transform(X_train)
X_test_tfidf = vectorizer.transform(X_test)
```

```
In [7]: nb_model = MultinomialNB()
nb_model.fit(X_train_tfidf, y_train)
y_pred_nb = nb_model.predict(X_test_tfidf)
print("\n===== Naive Bayes =====")
print("Accuracy:", accuracy_score(y_test, y_pred_nb))
print(classification_report(y_test, y_pred_nb))
```

```
===== Naive Bayes =====
Accuracy: 0.9668161434977578
```

	precision	recall	f1-score	support
0	0.96	1.00	0.98	965
1	1.00	0.75	0.86	150
accuracy			0.97	1115
macro avg	0.98	0.88	0.92	1115
weighted avg	0.97	0.97	0.96	1115

```
In [8]: lr_model = LogisticRegression(max_iter=1000)
lr_model.fit(X_train_tfidf, y_train)
y_pred_lr = lr_model.predict(X_test_tfidf)
print("\n===== Logistic Regression =====")
print("Accuracy:", accuracy_score(y_test, y_pred_lr))
print(classification_report(y_test, y_pred_lr))
```

```
===== Logistic Regression =====
Accuracy: 0.9524663677130045
```

	precision	recall	f1-score	support
0	0.95	1.00	0.97	965
1	0.97	0.67	0.79	150
accuracy			0.95	1115
macro avg	0.96	0.83	0.88	1115
weighted avg	0.95	0.95	0.95	1115

```
In [9]: svm_model = LinearSVC()
svm_model.fit(X_train_tfidf, y_train)
y_pred_svm = svm_model.predict(X_test_tfidf)
print("\n===== Support Vector Machine (SVM) =====")
```

```
print("Accuracy:", accuracy_score(y_test, y_pred_svm))
print(classification_report(y_test, y_pred_svm))
```

==== Support Vector Machine (SVM) ====

Accuracy: 0.97847533632287

	precision	recall	f1-score	support
0	0.98	0.99	0.99	965
1	0.96	0.87	0.92	150
accuracy			0.98	1115
macro avg	0.97	0.93	0.95	1115
weighted avg	0.98	0.98	0.98	1115

c:\Users\damma\anaconda4\envs\tf\lib\site-packages\sklearn\svm\\_classes.py:32: FutureWarning: The default value of `dual` will change from `True` to `auto` in 1.5. Set the value of `dual` explicitly to suppress the warning.  
warnings.warn(

```
In [12]: # Simple prediction
msg = "Congratulations! You've won a $1000 gift card."
msg_tfidf = vectorizer.transform([msg])
pred = svm_model.predict(msg_tfidf)[0]
print("Spam" if pred == 1 else "Ham")
```

Spam